

# High Performance Computing – Tutorial

Dr.-Ing. Martin Bernreuther

Höchstleistungsrechenzentrum der  
Universität Stuttgart



SS 2020

# Organizational

- for the creation of a VIS-SgS-Pool account, the following data is needed:
  - surname, first name, (gender)
  - matriculation number
  - login name GS/HS-Pool (if existant)
  - E-Mail
  - major subject, terms of studying
- E-Mail to [bernreuther@hlrs.de](mailto:bernreuther@hlrs.de) (ASAP)
- IPVS/SgS-Cluster: submit scan of completed, signed form [IPVS.antrag\\_rechner.pdf](#)
- bwUniCluster account: see <http://www.hlrs.de/solutions-services/academic-users/bwunicluster-access/>  
(mailto:bwunicluster@hlrs.de referencing this course)

## VIS/IPVS remote login

- Login using ssh, e.g.  
ssh -X \$USER@ssh1.visus.uni-stuttgart.de  
ssh -X \$USER@ipvslogin.informatik.uni-stuttgart.de
- further connect to a VISSGS pool machine, e.g.  
ssh -X visgs01  
check, if other users are logged in with e.g. who or w
- you might use e.g. scp to transfer data (or applications like FileZilla . . .)
- consider using screen or tmux (or tmate), to be able to reconnect after loosing connection or to share the terminal.

Sharing a session with screen:

- use -S <SESSIONNAME> option to name the session
  - Ctrl-a :multiuser on, Ctrl-a :acladd OTHERUSER
  - OTHERUSER on the same machine might attach with screen -x FIRSTUSER/SESSIONNAME
- to access a graphical desktop remotely, you might use e.g. x2go

## Environment modules

- to adjust the environment for a software (setting or expanding environment variables, e.g. add a path to the PATH variable), clusters usually use environment modules
- `module avail` shows, which modules are available
- `module show <module>` shows changes for *module*
- `module load <module>` executes the changes for *module* (e.g. `module load gcc` or `module load gcc/gcc-7.1`)
- `module unload <module>` revokes the changes for *module*
- `module list` list active modules
- module will check the paths listed in the MODULEPATH variable for modules
- `$HOME/.modulerc` to set up user specific module changes for each login  
(e.g. `module load use.own` adds `$HOME/privatemodules` to the MODULEPATH)

# Hardware properties

- Linux shows information about the CPUs in a virtual text file: `/proc/cpuinfo`  
the capabilities of the host cpu can be determined e.g. with  
`grep '^flags' /proc/cpuinfo | uniq -c`  
There are also commands like `lscpu`, `inxi` or `hwinfo`  
(`nproc` shows number of PEs)
- there's information about the memory at `/proc/meminfo`  
and there are commands like e.g. `free`, `vmstat`,...
- the cache structure and sizes can be determined with e.g.  
`likwid --topology -g`, or within a bash

```
for d in /sys/devices/system/cpu/cpu*/cache/index*; do
  echo -n " 'cat_$d/type ' 'cat_$d/level '"
  echo -n " , 'cat_$d/ways_of_associativity '-way"
  echo -n " , _cl_ 'cat_$d/coherency_line_size '"
  echo "      " , 'cat_$d/size '"
done | sort | uniq -c
```

- also check `top` or `htop` before starting a benchmark

# Process/Thread pinning

avoid cpu-migrations. . .

## ■ taskset

- `taskset <cpu_mask> <cmd_with_options>`,  
e.g. `<cpu_mask> 0x1` to use core 0 or  
`0x3ff` to use the cores 0-9
- `taskset -a -p <PID>`,  
to retrieve the `cpu_mask` of all tasks of a running process  
`taskset -p <cpu_mask> <PID>`,  
to change the process `cpu_mask`

## ■ numactl

- `numactl --hardware` and `numactl --show` to display information
- `numactl --physcpubind=0,1,2,3 <cmd_with_options>`  
to use core 0-3
- `numactl --cpunodebind=0 --membind=0 <cmd_with_options>`

## ■ likwid-pin

## ■ also see OpenMP and MPI

# C(++) time measurement

- **processor time** (cycles spent for the program)
  - **clock**
- **walltime** (for parallel programs usually the walltime is measured)
  - **time.h**: `time`
  - **time.h**: `clock_gettime`
  - **sys/time.h**: `gettimeofday`
  - **omp.h**: `omp_get_wtime` (OpenMP)
  - **mpi.h**: `MPI_Wtime` (MPI)
  - **chrono**: `std::chrono::system_clock::now` or  
`std::chrono::high_resolution_clock::now`

# Part I

## Tutorial: Streaming benchmark



# Stream benchmark

- Create a C-program to calculate vector operations like

$$1 \quad a[i] = b[i] \quad \forall i = 0, \dots, N-1$$

$$2 \quad a[i] = b[i] + c[i] \quad \forall i = 0, \dots, N-1$$

$$3 \quad a[i] = s \cdot b[i] + c[i] \quad \forall i = 0, \dots, N-1$$

$$4 \quad a[i] = b[i] \cdot c[i] + d[i] \quad \forall i = 0, \dots, N-1$$

also see <http://www.cs.virginia.edu/stream/> or likwid-bench

→ we'll focus on the last variant, the Schönauer vector triade

- Vectors are one-dimensional dynamical arrays of length  $N$ , which initially are allocated and initialized with values.  
(thus there's already a “first touch” and cache read).
- the (BLAS1) vector operations will be executed  $nrepeat$  times, the walltime is measured with e.g. `gettimeofday` to calculate an average.
- determine the average FLOp/s and memory bandwidth for varying  $N$  (and *STRIDE*, double vs. single precision. . .) and visualize it with e.g. `gnuplot`
- these stream benchmarks are memory-bound!

# valgrind

- valgrind takes control of a program and runs it on a synthetic CPU (The program binary needs not to be modified in any way.)
- valgrind supports several platforms, like currently (v3.15) e.g.
  - x86/Linux: up to and including SSSE3, but not higher – no SSE4, AVX, AVX2. This target is in maintenance mode now..
  - AMD64/Linux: up to and including AVX2. This is the primary development target and tends to be well supported.
  - ARM/Linux: supported since ARMv7.
  - ARM64/Linux: supported for ARMv8.
  - ...
- before execution of a code part it is handed to a tool (formerly referred to as skin), which adds instrumentation code

another similar utility: intel pin

## valgrind tools

- the cache behaviour of the benchmarks can be simulated using the cachegrind tool: `valgrind --tool=cachegrind <binary>`
- especially for applications with multiple functions, the “callgrind” tool can determine these performance values and more on a functional level and visualize it with “kcachegrind”
- tools for multithreading applications:  
(e.g. pthreads; there might be problems with (Linux) futex/OpenMP)
  - “helgrind” – a thread error detector
  - “drd” (**D**ata **R**ace **D**etection) – a thread error detector
- tools for memory usage:
  - “memcheck” – a memory error detector (default)
  - “massif” – a heap profiler
  - “DHAT” (**D**ynamic **H**eap **A**nalysis **T**ool)

# Using cachegrind & callgrind

compile with debugging information (“-g”) for source file related output

## ■ Cachegrind

- `valgrind --tool=cachegrind --branch-sim=yes ./stream_ref_func 10000`  
produces a text file (`cachegrind.out.*`) with cache usage information
- `cg_annotate --auto=yes --show-percs=yes cachegrind.out.*`  
to show the cache ref./misses for each line of the source code

## ■ Callgrind (extension of cachegrind)

- `valgrind --tool=callgrind --cache-sim=yes ./stream_ref_func 1000`  
produces a text file `callgrind.out.*`  
you might also use more options, like e.g.  
`--simulate-wb=yes --simulate-hwpref=yes --cacheuse=yes`
- `kcachegrind callgrind.out.*`  
to visualize the data



remember: cache accesses refer to cachelines

- **Performance Counters for Linux**/perf is a userspace performance analyzing tool of the Linux kernel  
(unlike e.g. for OProfile no daemon is required; reduce /proc/sys/kernel/perf\_event\_paranoid)
- list events with e.g. `perf list --long-desc --details`  
(`perf list hw pmu` to list hardware events only)



```
perf top [-e <event_type>]
perf stat -ad -- <cmd_with_options>
perf record <cmd_with_options>|-p <PID>
perf report
perf annotate -d <cmd>
```



# Linux perf stat example

## ■ perf stat example output:

```
$ export LC_NUMERIC=C
$ perf stat -d taskset 1 ./stream_ref.gcc 100000 1
Performance counter stats for 'taskset 1 ./stream_ref.gcc 100000 1':
    2.190340      task-clock (msec)          #    0.864 CPUs utilized
           0      context-switches        #    0.000 K/sec
           0      cpu-migrations          #    0.000 K/sec
          920      page-faults            #    0.420 M/sec
       7504129      cycles                  #    3.426 GHz
       7072664      instructions             #    0.94 insn per cycle
       1287923      branches                # 588.001 M/sec
          21833      branch-misses          #   1.70% of all branches
       1249007      L1-dcache-loads         # 570.234 M/sec
          150852      L1-dcache-load-misses     #  12.08% of all L1-dcache hits
<not counted>      LLC-loads                    (0.00%)
<not counted>      LLC-load-misses             (0.00%)
    0.002535691 seconds time elapsed

$ perf stat -e cycles,instructions,\
  fp_arith_inst_retired.scalar_double,fp_arith_inst_retired.256b_packed_double\
  ./stream_avx.gcc 100
Performance counter stats for './stream_ref-avx2fma3.gcc 100':
    1628452      cycles
    897762      instructions              #    0.55 insn per cycle
           6      fp_arith_inst_retired.scalar_double
          50      fp_arith_inst_retired.256b_packed_double
    0.000947457 seconds time elapsed

(scalar FP instructions used for time calculation; FMA instructions count twice)
```

# Performance Application Programming Interface

- The **P**erformance **A**pplication **P**rogramming Interface (PAPI) is a machine independent API to provide access to performance counters
- check PAPI events availability and get information with `papi_avail`
- include `papi.h` and compile/link with `–lpapi`
- might use e.g. the high level function  
`int PAPI_flops( float *rtime, float *ptime, long long *fpops, float *mflops);`  
before and after the loop(s) to determine the FLOP/s...

(How are e.g. AVX instructions counted?)

## Loop unrolling for the vector triad

```
for ( i=0; i<N; i+=SIMD_WIDTH)
{ /* for SIMD_WIDTH==4 */
  a [ i ] = b [ i ] * c [ i ] + d [ i ];
  a [ i+1 ] = b [ i+1 ] * c [ i+1 ] + d [ i+1 ];
  a [ i+2 ] = b [ i+2 ] * c [ i+2 ] + d [ i+2 ];
  a [ i+3 ] = b [ i+3 ] * c [ i+3 ] + d [ i+3 ];
  /* might be replaced by a single SIMD instr. */
}
```

- with explicit successive calculations of consecutive loop iterations, it might be easier for the compiler to optimize and generate AVX-instructions.

It's possible to check the compiler dependent optimization looking at an assembler listing. [objdump](#)

- shouldn't be necessary for this simple example, but sometimes still does not lead to the hoped for result...



# AVX intrinsics variant


- explicit utilization of AVX intrinsics to increase performance
- 4 double precision values (with 8 Byte = 64 Bit each) are processed simultaneously in 32 Byte (= 256 Bit) registers.

Assumption:  $N$  is divisible by 4 ( $N\%4 == 0$ )

(otherwise “padding” with dummy-elements or special treatment with scalar operations)

- Data-alignment
  - “packed double” operations access data with starting addresses being  $4 \cdot 8 = 32$  byte aligned.
  - to ensure alignment the memory allocation should use suitable functions to replace malloc.

# Memory allocation with data alignment



- “MultiMedia” Intrinsics aligned memory allocation:

```
#include <mm_malloc.h>
void* _mm_malloc (int size, int align);
void _mm_free (void *p);
```

- POSIX posix\_memalign, C11 aligned\_alloc :

```
#include <stdlib.h>
int posix_memalign(void **memptr, size_t alignment, size_t size);
void *aligned_alloc(size_t alignment, size_t size);
```


- gcc offers e.g. keyword `__attribute__((aligned(X)))` for static arrays and structs/classes
- C++11: `alignas(X)`; also check `std::align` and `std::aligned_storage`
- Library functions: `Boost.Align`

# AVX intrinsics version details

## ■ AVX-Commands

- to use AVX commands it's necessary to  
**#include** <immintrin.h>
- usage of AVX double precision registers: `__m256d va;`
- set values: `va=_mm256_set_pd(a3,a2,a1,a0);` or `va=_mm256_set1_pd(a);`
- load values (aligned) from an array: (packed double)  
`va=_mm256_load_pd(&a[i]);`
- write values (aligned) to an array: `_mm256_store_pd(&a[i],va);`  
alternative: `_mm256_stream_pd(&a[i],va);`
- multiplication, addition:  
`va=_mm256_mul_pd(vb,vc);`  
`va=_mm256_add_pd(va,vd);`  
`//va=_mm256_fmadd_pd(vb,vc,vd);`
- AMD Interlagos, Intel Haswell:  
also HW-supported **F**used **M**ultiply **A**dd 4

- compiler optimization also applies SIMD-instructions  
(possibly specify architecture, e.g. `gcc -mavx`)



## AVX intrinsics variant: vector triad excerpt

```
for ( i=0; i<N; i+=SIMD_WIDTH) {  
    vb=_mm256_load_pd(&b[ i ] );  
    vc=_mm256_load_pd(&c[ i ] );  
    vd=_mm256_load_pd(&d[ i ] );  
#ifdef USE_AVX_FMA4  
    va=_mm256_fmadd_pd( vb , vc , vd );  
#else  
    va=_mm256_add_pd( _mm256_mul_pd( vb , vc ) , vd );  
#endif  
#ifdef USE_AVX_STREAM  
    _mm256_stream_pd(&a[ i ] , va );  
#else  
    _mm256_store_pd(&a[ i ] , va );  
#endif  
}
```

## more AVX intrinsics variations

- replace `_mm256_load_pd(&a[i])` with `_mm256_loadu_pd(&a[i])` and `_mm256_store_pd(&a[i])` with `_mm256_storeu_pd(&a[i])` to allow unaligned memory access and compare the performance.
- replace `_mm256_load_pd(&a[i])` with `va=_mm256_set_pd(a[i+3],a[i+2],a[i+1],a[i]);` to simulate gathering the data from different variables instead of loading aligned array data.
- setting a mask with `_mm256_set_epi64x` and masking loads with `_mm256_maskload_pd` and stores with `_mm256_maskstore_pd` will add flexibility for e.g. strided access or conditional statements. Using

```
__m256i vmask=_mm256_set_epi64x(-1,-1,-1,-1);
```

will retain the functionality of the unmasked version, but probably still influence the performance.

## SSE2 intrinsics variant

- as a comparison and also for CPUs, which do not support AVX extensions (cmp. “flags” in “/proc/cpuinfo”), analogous also a SSE2 variant can be created
- SSE: 2 double precision values are processed simultaneously.
- alignment:  $2 \cdot 8 = 16$  byte
- **#include** <emmintrin.h> (for SSE2)  
MMX:mmmintrin.h, SSE:xmmmintrin.h, SSE3:pmmintrin.h, SSSE3:tmmintrin.h, SSE4.1:smmintrin.h, SSE4.2:nmmintrin.h, SSE4A:ammintrin.h, . . .
- SSE2 double precision registers: `__m128d` va;
- prefix for the commands: “`_mm_`”
- set values: `va=_mm_set_pd(a1,a0);`
- no “fma” (fused multiply-add) or “stream” instruction
- `_mm_prefetch(ADDRESS,LOCALITYHINT)` might give hints for prefetching

# Intel intrinsics: performance

Performance (Latency; Throughput [CPI]):

Instruction	Architecture			
	Ivy Bridge	Haswell	Broadwell	Skylake
_mm_load_pd	1; 1	1; 0.5	1; 0.5	1; 0.33
_mm_store_pd	1; 1	1; 0.5	1; 0.5	1; 0.33
_mm_add_pd	3; 1	3; 1	3; 1	4; 0.5
_mm_mul_pd	5; 1	5; 0.5	5; 0.5	3; 0.5
_mm256_load_pd	1; 1	1; 0.5	1; 0.5	1; 0.25
_mm256_store_pd	1; 1	1; 0.5	1; 0.5	1; 0.25
_mm256_stream_pd	-; 1	-; 1	-; 1	-; 1
_mm256_add_pd	3; 1	3; 1	3; 1	4; 0.5
_mm256_mul_pd	5; 1	5; 0.5	3; 0.5	4; 0.5
_mm256_fmadd_pd	-	5; 0.5	5; 0.5	4; 0.5
_mm256_loadu_pd	1; 1	1; 0.5	1; 0.5	1; 0.25
_mm256_storeu_pd	1; 1	1; 0.5	1; 0.5	1; 0.25


(source: Intel Intrinsics Guide)

## AVX512 intrinsics variant

- for actual (Intel) CPUs analogous also a AVX512 variant can be created (check “flags” in “/proc/cpuinfo”)
- AVX512: 8 double precision values are processed simultaneously.
- alignment:  $8 \cdot 8 = 64$  byte
- **#include** <immintrin.h> (like AVX and AVX2)
- AVX512 double precision registers: `__m512d va;`
- prefix for the commands: “`_mm512_`”
- set values: `va=_mm512_set_pd(a7,a6,a5,a4,a3,a2,a1,a0);`



# Intel CPUs: AVX512 levels & support



source:instlatx64.atw.hu

## Excurs: Approximation of $\exp()$ function

- a Taylor series can be used to approximate and replace a function  $f(x)$ :  $f(x) = \lim_{o \rightarrow \infty} \sum_{i=0}^o \frac{f^{(i)}(x_0)}{i!} (x - x_0)^i$
- to approximate  $\exp(x) = e^x$  with this polynomial, we get  $\exp(x) \approx \sum_{i=0}^o \frac{\exp(x_0)}{i!} (x - x_0)^i$  and for the Maclaurin series ( $x_0 = 0$ )  $\exp(x) \approx \sum_{i=0}^o \frac{1}{i!} x^i$
- $o$  can be reduced using multiple  $x_0$  (e.g. based on the exponent of  $x$ )
- the sum can be calculated using FMA within a loop  
 $sum += a_i \cdot \tilde{x}^i$  (Horner's method:  $sum = sum \cdot \tilde{x} + a_{o-i}$ ),  
 where the factors  $a_i$  can also be pre-computed.
- SIMD units of width  $w$  can be used with loop unrolling and calculating partial sums first.  
 (A  $O(\log(w))$  algorithm to reduce the final sum of  $w$  partial sums might be inefficient for small  $w$ .)
- use of SIMD calculating  $\exp(x_i)$  for each element of a vector?


## ISPC variant

- as a starting point, a stream version is used, where the kernel is moved to a (external) function.
- to create a Intel SPMD Program Compiler version, this function is replaced with a ISPC version compiled with the ispc compiler:

```
export void stream(uniform const Tindex N
                  ,uniform Tfloat a[]
                  ,uniform const Tfloat b[]
                  ,uniform const Tfloat c[]
                  ,uniform const Tfloat d[]
                  ) {
    foreach (i = 0 ... N) {
        a[i]=b[i]*c[i]+d[i];
    }
}
```


# vTriad Stream Benchmark on a VIS-SgS-Computer

Schönauer Triad Stream benchmark, visgs (SandyBridge)



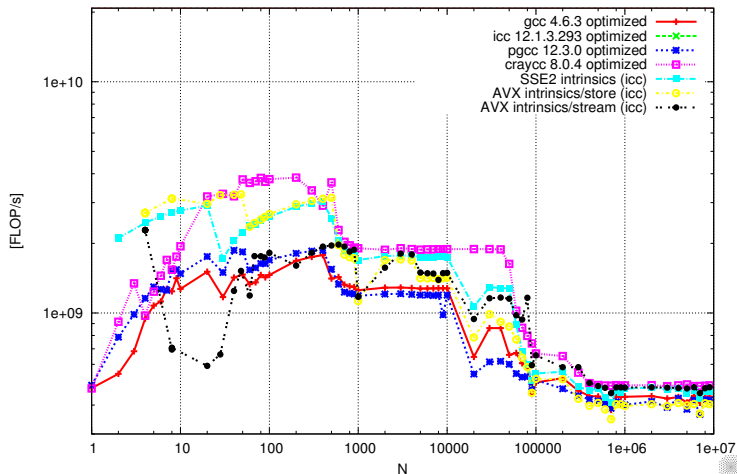
# vTriad Stream Benchmark on a VIS-SgS-Computer

Schönauer vector Triad Stream benchmark on visgs01 (Intel i5-9600K, gcc9.3.1)



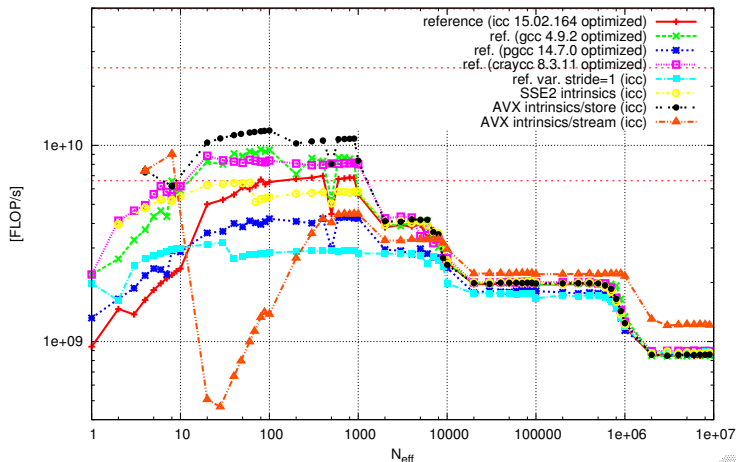
## vTriad Stream Benchmark on Hermit

Schönaauer Triad Stream benchmark, hermit (Interlagos)



# vTriad Stream Benchmark on Hornet/Hazelhen

Schönauer Triad Stream benchmark, hornet (Haswell)



# objdump

- with the help of e.g. `objdump -d <objfile>` it's possible to take a look at the assembler code of the variants.
- setting the `-M intel-mnemonic` option, objdump will produce Intel instead of AT&T mnemonic code syntax.
- alternatively the gcc, clang, icc and pgcc can generate an assembler listing using the `-S` option
- the main program can be found in the “<main>” section, the computation part between the “gettimeofday” calls



# x86\_64 assembler

- a command consists of the

- 1 instruction (Opcode)

- 2 up to 2 operands

- Register

- **General Purpose Registers (64-Bit):**

- RAX (Accumulator), RBX (Base register), RCX (Counter), RDX (Data register), RBP (Base-Pointer), RSI (Source-Index), RDI (Destination-Index), RSP (Stack-Pointer), R8, . . . , R15, . . .

- additional 64-bit registers:

- RIP (Instruction-Pointer), RFLAGS (FLAGS, status), XMM0, . . . , XMM15 (SSE), YMM0, . . . , YMM15 (AVX), ZMM0, . . . , ZMM31 (AVX512)

- e.g. for 64-bit registers rax, r8:

- eax, r8d address lowest 32 bits; ax, r8w lowest 16 bits;

- al, r8b lowest 8 bits (ah higher 8 bits of ax)

# x86\_64 assembler

- a command consists of the
  - 1 instruction (Opcode)
  - 2 up to 2 operands
- **Registernames**  
 RAX,RBX,RCX,RDX,RBP,RSI,RDI,RSP,R8,...,R15,...,XMM0,...,XMM15,YMM0,...,YMM15,...
- assembly source file sections:  
 Data, Block Started by Symbol (uninitialized data), Text (code)
- memory addressing (for  $a[i]$  or  $*(a+i)$ ):  
 $Base + Displacement + i * Size$
- **packed double** suffix: “vaddpd” and “vmulpd” e.g. denote addition- and multiplication SSE2/AVX instructions.
- can these AVX-instructions already be found in the standard reference variant? (using compiler optimization-options)

## objdump: stream\_ref vTriad computation part

```

4007f0:    e8 0b fe ff ff          call    400600 <gettimeofday@plt>
4007f5:    31 c0                   xor     eax,eax
4007f7:    31 d2                   xor     edx,edx
4007f9:    48 8b 4d 90             mov     rcx,QWORD PTR [rbp-0x70]

4007fd:    f2 41 0f 10 04 d7       movsd   xmm0,QWORD PTR [r15+rdx*8]
400803:    f2 0f 59 04 d1          mulsd   xmm0,QWORD PTR [rcx+rdx*8]
400808:    48 8b 4d 88             mov     rcx,QWORD PTR [rbp-0x78]
40080c:    f2 0f 58 04 d1          addsd   xmm0,QWORD PTR [rcx+rdx*8]
400811:    f2 41 0f 11 04 d6       movsd   QWORD PTR [r14+rdx*8],xmm0
400817:    48 ff c2               inc     rdx
40081a:    4c 39 ea               cmp     rdx,r13
40081d:    75 da               jne     4007f9 <main+0x1a9>
40081f:    48 ff c0               inc     rax
400822:    49 39 c4               cmp     r12,rax
400825:    75 d0               jne     4007f7 <main+0x1a7>
400827:    48 8d 7d c0             lea     rdi,[rbp-0x40]
40082b:    31 f6                   xor     esi,esi
40082d:    48 89 55 80             mov     QWORD PTR [rbp-0x80],rdx

400831:    e8 ca fd ff ff          call    400600 <gettimeofday@plt>

```



## objdump: stream\_sse2 vTriad computation part

```

40073e:    e8 cd fe ff ff          call    400610 <gettimeofday@plt>
400743:    31 c9                   xor     ecx,ecx
400745:    4c 89 e2                mov     rdx,r12
400748:    0f 1f 84 00 00 00 00    nop    DWORD PTR [rax+rax*1+0x0]
40074f:    00                      nop
400750:    48 8b 75 98             mov     rsi,QWORD PTR [rbp-0x68]
400754:    31 c0                   xor     eax,eax
400756:    66 2e 0f 1f 84 00 00    nop    WORD PTR cs:[rax+rax*1+0x0]
40075d:    00 00 00                nop
400760:    c4 c1 79 28 04 c7      vmovapd xmm0,XMMWORD PTR [r15+rax*8]
400766:    c5 f9 59 04 c6          vmulpd  xmm0,xmm0,XMMWORD PTR [rsi+rax*8]
40076b:    c4 c1 79 58 04 c6      vaddpd  xmm0,xmm0,XMMWORD PTR [r14+rax*8]
400771:    c5 f8 29 04 c3          vmovaps XMMWORD PTR [rbx+rax*8],xmm0
400776:    48 83 c0 02             add     rax,0x2
40077a:    48 39 d0                cmp     rax,rdx
40077d:    75 e1                   jne     400760 <main+0x100>
40077f:    48 83 c1 01             add     rcx,0x1
400783:    49 39 cd                cmp     r13,rcx
400786:    75 c8                   jne     400750 <main+0xf0>
400788:    48 8d 7d c0             lea     rdi,[rbp-0x40]
40078c:    31 f6                   xor     esi,esi

40078e:    e8 7d fe ff ff          call    400610 <gettimeofday@plt>

```

## objdump: stream\_avx vTriad computation part

```

40074d:    e8 be fe ff ff          call    400610 <gettimeofday@plt>
400752:    4c 8b 5d 90             mov     r11,QWORD PTR [rbp-0x70]
400756:    31 c9                   xor     ecx,ecx
400758:    4c 89 ea             mov     rdx,r13
40075b:    0f 1f 44 00 00         nop     DWORD PTR [rax+rax*1+0x0]

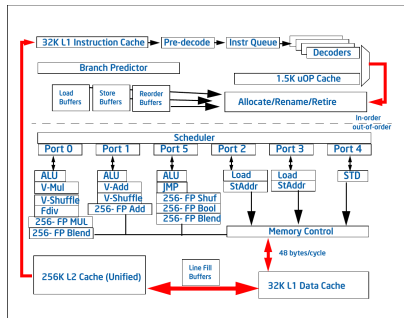
400760:    31 c0                   xor     eax,eax
400762:    66 0f 1f 44 00 00     nop     WORD PTR [rax+rax*1+0x0]

400768:    c4 c1 7d 28 04 c7     vmovapd ymm0,YMMWORD PTR [r15+rax*8]
40076e:    c4 c1 7d 59 04 c4     vmulpd  ymm0,ymm0,YMMWORD PTR [r12+rax*8]
400774:    c4 c1 7d 58 04 c3     vaddpd  ymm0,ymm0,YMMWORD PTR [r11+rax*8]
40077a:    c5 fd 29 04 c3     vmovapd YMMWORD PTR [rbx+rax*8],ymm0
40077f:    48 83 c0 04         add     rax,0x4
400783:    48 39 d0             cmp     rax,rdx
400786:    75 e0             jne     400768 <main+0x108>
400788:    48 83 c1 01         add     rcx,0x1
40078c:    49 39 ce             cmp     r14,rcx
40078f:    75 cf             jne     400760 <main+0x100>
400791:    48 8d 7d c0         lea     rdi,[rbp-0x40]
400795:    31 f6             xor     esi,esi
400797:    4c 89 5d 88             mov     QWORD PTR [rbp-0x78],r11
40079b:    c5 f8 77             vzeroupper
40079e:    e8 6d fe ff ff          call    400610 <gettimeofday@plt>

```

# Intel Architecture Code Analyzer

- the Intel Architecture Code Analyzer (IACA) perform a static analysis of the data dependencies, the throughput and the latency of binary code snippets.
- in the source code parts can be labeled through markers  
(instrumentation of the source code)
- IACA shows the assignment of kernel instructions to processor ports and determine the critical path



Intel® 64 and IA-32  
Architectures Optimization Reference Manual,  
S. 2-4 (42), Figure 2-1

## IACA: stream\_avx vTriad comp. part (gcc -O3; intel Sandy Bridge)

```
$ iaca -64 -arch SNB stream_avx.iaca
```

```
Intel(R) Architecture Code Analyzer Version - 2.2 build:1aef335 (Wed, 28 Dec 2016 15:14:25 +0200)
```

```
Analyzed File - stream_avx.iaca
```

```
Binary Format - 64Bit
```

```
Architecture - SNB
```

```
Analysis Type - Throughput
```

## Throughput Analysis Report

Block Throughput: 3.00 Cycles

Throughput Bottleneck: Port2\_DATA, Port3\_DATA

Port Binding In Cycles Per Iteration:

Port	0 - DV	1	2 - D	3 - D	4	5
Cycles	1.0 0.0	1.0 2.0	3.0 2.0	3.0 2.0	2.0 2.0	

N - port number or number of cycles resource conflict caused delay, DV - Divider pipe (on port 0)

D - Data fetch pipe (on ports 2 and 3), CP - on a critical path

F - Macro Fusion with the previous instruction occurred

\* - instruction micro-ops not bound to a port

^ - Micro Fusion happened

# - ESP Tracking sync uop was issued

@ - SSE instruction followed an AVX256/AVX512 instruction, dozens of cycles penalty is expected

X - instruction not supported, was not accounted in Analysis

Num Of		Ports pressure in cycles							
Uops	0 - DV	1	2 - D	3 - D	4	5			
1			0.5	1.0	0.5	1.0			CP
2	1.0		0.5	1.0	0.5	1.0			CP
2		1.0	0.5	1.0	0.5	1.0			CP
2			0.5				2.0		
1								1.0	
1								1.0	
0F									

Total Num Of Uops: 9



```
CP vmovapd ymm0, ymmword ptr [r15+rax*8]
CP vmulpd ymm0, ymm0, ymmword ptr [r12+rax*8]
CP vaddpd ymm0, ymm0, ymmword ptr [r11+rax*8]
vmovapd ymmword ptr [rbx+rax*8], ymm0
add rax, 0x4
cmp rax, rdx
jnz 0xffffffffffffffffda
```





# OSACA, LLVM-MCA

- alternatives to IACA (reaching EOL April 2019):
  - **Open Source Architecture Code Analyzer**  
(using kerncraft)
  - **Low Level Virtual Machine - Machine Code Analyzer**
  - also pmu-tools

## ARM variant

- compile under specification of the architecture (on “CubieTower”):

```
gcc -4.7 -march=armv7 -mfpu=neon -O3 \  
      -funsafe-math-optimizations -lm \  
      stream_ref_float.c -o stream_ref_float
```

- check CPU frequency:

```
$ dmesg | grep freq  
$ cd /sys/devices/system/cpu/cpu0/cpufreq  
$ cat scaling_available_governors scaling_governor  
$ echo 'performance' > scaling_governor  
$ cat cpuinfo_cur_freq  
$ echo 1008000 > scaling_setspeed
```

# ARM NEON intrinsics variant (single precision)

- header file: **#include** <arm\_neon.h>


## SSE2

## NEON

<code>_mm128 v;</code>	<code>float32x4_t v;</code>
<code>v=_mm_set_ps(s,s,s,s);</code>	<code>v=vdupq_n_f32(s)</code>
<code>v=_mm_set_ps(s0,s1,s2,s3);</code>	<code>v=vsetq_lane_f32(s0,v,0);</code>
	<code>v=vsetq_lane_f32(s1,v,1);</code>
	<code>v=vsetq_lane_f32(s2,v,2);</code>
	<code>v=vsetq_lane_f32(s3,v,3);</code>
<code>v=_mm_load_ps(p);</code>	<code>v=vld1q_f32(p);</code>
<code>_mm_store_ps(p,v);</code>	<code>vst1q_f32(p,v);</code>
<code>v=_mm_add_ps(va,vb);</code>	<code>v=vaddq_f32(va,vb);</code>
<code>v=_mm_mul_ps(va,vb);</code>	<code>v=vmulq_f32(va,vb);</code>

(s. e.g. ARM® NEON™ Intrinsic Reference)

# vTriad Stream Benchmark on a Cubietruck (SP)



## PThread variant (1)

- introducing an array with elements of a structure, which contain the necessary data for each thread, e.g.

```
typedef struct {  
    unsigned int p, numthreads; /* ThreadID, Number of threads */  
    unsigned int N;             /* Number of elements */  
    double *a,*b,*c,*d;         /* arrays */  
} t_thread_work;
```

and in the main program

```
t_thread_work thread_work[NUMTHREADS];  
for (p=0; p<NUMTHREADS; ++p) {  
    thread_work[p].N=N;  
    thread_work[p].p=p;  
    thread_work[p].numthreads=NUMTHREADS;  
    thread_work[p].a=a; thread_work[p].b=b;  
    thread_work[p].c=c; thread_work[p].d=d;  
}
```

## PThread variant (2)

- introducing an array with elements of a structure, which contain the necessary data for each thread
- function with a pointer to such a structure as parameter, which contains the core functionality, e.g.

```
void thread_stream(t_thread_work *w) {  
    /* set workpackage for each thread separately ,  
       dependent on p,... */  
    for( i=i1 ; i<i2 ; i+=di )  
        w->a [ i ] = w->b [ i ] * w->c [ i ] + w->d [ i ] ;  
}
```

## PThread variant (3)

- introducing an array with elements of a structure, which contain the necessary data for each thread
- function with a pointer to such a structure as parameter, which contains the core functionality
- creating new threads, to let them execute this function  
threads will join again afterwards → barrier

```
for (p=0;p<NUMTHREADS;++p) {
    pthread_create(&thread[p],NULL
                  ,(void*)(*(void*))&thread_stream,(void*)&(thread_work[p]));
}
for (p=0;p<NUMTHREADS;++p) pthread_join(thread[p],NULL);
```

include the function & datatype definitions before through  
**#include** <pthread.h>

## PThread variant (4)

- introducing an array with elements of a structure, which contain the necessary data for each thread
- function with a pointer to such a structure as parameter, which contains the core functionality
- creating new threads, to let them execute this function

```
int pthread_create( pthread_t *thread, NULL  
                  , void *(*start_routine)(void*), void *arg );
```

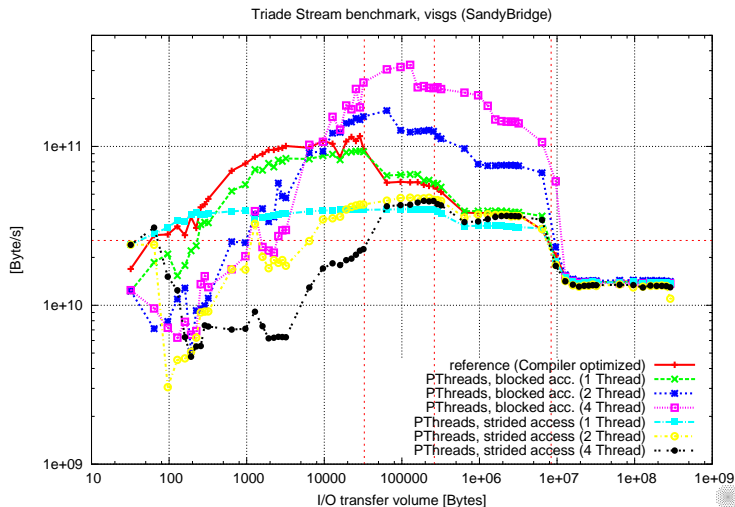
threads will join again afterwards → barrier

```
int pthread_join( pthread_t thread, NULL );
```

- translate with necessary options (e.g. `gcc -pthread`)
- test run with `valgrind --tool=helgrind` (or `--tool=drd`)
- PThreads can be combined with SIMD-variants
- optional: C++11 (or C++17) and JavaThreads variants (runtime comparison)




# PThread vTriad Stream Benchmark @VIS-SgS



## OpenMP variant

- initial OpenMP version starting from the reference-C-variant:
  - the most inner vector based loop is parallelized:

**#pragma omp parallel for**




- utilization of `omp_get_wtime()` to measure the time  
the usage of library functions requires to include the header file:

```
#ifdef _OPENMP
#include <omp.h>
#endif
```

- compile with necessary options (e.g. `gcc -fopenmp`)
- set environment variable `OMP_NUM_THREADS`

## OpenMP variant (2)

- first modification:
  - parallel block including repetition loop:
   
**#pragma omp parallel**
  
(Which variables have to be declared as “private”? Synchronization? Barriers required?)
  - inner  $i$ -loop is executed collectively
   
**#pragma omp for**



- Impact of the “schedule” directive to the runtime?
- “if” directive for the parallel block?

## OpenMP variant (2)

- first modification: alternative version

- parallel block including repetition loop:

**#pragma omp parallel**

(Which variables have to be declared as "private"? Synchronization? Barriers required?)


- inner  $i$ -loop is executed collectively

- each thread executes only a part of the loop iterations:

**if** ( $i \% \text{ompnumthreads} == \text{ompthreadid}$ ) ... with

**int** ompnumthreads=omp\_get\_num\_threads();


**int** ompthreadid=omp\_get\_thread\_num();



- Change of the runtime? Reasons?

## OpenMP variant (3)

- second modification for NUMA systems (and to avoid “false share”) using **#pragma omp for** again (e.g. to benchmark shared L3 Cache):
  - Data partitioning  
(each thread initializes and computes its own vectors)
  - Expand parallel region to the initialization  
(alternatively another parallel region can be used)  
→ “first touch” by executing thread  
How do we measure the time now? → explicit barriers
  - output through a single thread (e.g. with **if (omp\_get\_thread\_num()==0)...**,  
**#pragma omp single**, **#pragma omp master** )
- Thread-pinning (to prevent the OS to migrate thread execution to other Cores)
  - with tools like “taskset”, “numactl”, “likwid-pin”, ... or
  - defined by environment variables
    - GOMP\_CPU\_AFFINITY (gcc), KMP\_AFFINITY (icc), ...
    - OMP\_PROC\_BIND, OMP\_PLACES (since OpenMP 4.0)



## OpenMP SIMD variant

- additional SIMD optimization possible. . .
- OpenMP also supports SIMD vectorization (since version 4)
- the most inner loop is vectorized with:  
**#pragma omp simd**
- Vectorization can be combined with thread parallelization:  
**#pragma omp parallel for simd**  
(or just **#pragma omp for simd** within a parallel region)
- How close can we get to the peak performance?  
Is there a saturation of the memory bus using multiple cores?

## Likwid-Bench

Likwid-Bench is part of the Likwid Performance Tools and a tool to perform streaming micro-benchmarks.

```
$ likwid-bench -a
```

will list the available benchmark kernels.

```
$ likwid-bench -l triad_avx
```

shows the properties of the `triad_avx` benchmark, and

```
$ likwid-bench -p
```

the available domains (CPU sockets, NUMA domains) Now e.g.

```
$ likwid-bench -t triad_avx -w S0:10kB
```

will e.g. run the Schönauer vector triad on the first socket (with one thread per core) for a working set size of 10 kB.